

- STARTING SHORTLY -

# A field guide to *Generative* *AI* for people who don't write code.

*In the next two hours you will prompt three different AI models, watch one of them confidently invent a fact, and leave with one tool you can use at work tomorrow morning.*

— TYPE IN CHAT

In one word — your current feeling about  
*artificial intelligence.*

Skeptical / Excited / Confused / Scared / Hopeful / Curious /  
Tired

**R**aise your hand   
in chat if AI has  
ever given you a  
confident,  
completely  
wrong answer.

— HOW WE'LL WORK

# For the next *two hours*.

---

I · CAMERAS

Encouraged but optional.

No call-outs for off-camera folks.

---

II · CHAT

Use it freely.

Questions, reactions, jokes — welcome.

---

III · MICS

Muted by default.

Unmute when called on or to share.

---

IV · QUESTIONS

Hot or parked.

Quick ones answered live. Bigger ones go to the parking lot doc.

---

*A real ten-minute break arrives at minute fifty-five.*

— VOTE IN THE POLL

# What will *ChatGPT* answer?

PROMPT

How many R's are in the word strawberry?

2

OPTION · A

3

OPTION · B

4

OPTION · C

*It'll be wrong*<sup>D</sup>

5<sup>/40</sup>

— CONTENTS

# Six topics. *Two hours.*

01

FOUNDATIONS

## Evolution of AI

Rules to foundation models.

02

LANDSCAPE

## Key models

GPT · DALL·E · Claude · Gemini.

03

MECHANICS

## How LLMs work

Tokens · context · temperature.

04

STACK

## The value chain

Models → platforms → apps.

05

RISK

## Limitations

Where AI breaks down.

06

CLIMAX

## Chatbot ceiling

Why single calls fall short.

6/40

— I. FOUNDATIONS —

# What is *Generative AI*?

*Evolution from rule-based systems to foundation models.*

— LIVE DEMO

The same Hindi phrase. *Two translations.*

PROMPT

मेरा दिल बाग बाग हो गया।

GOOGLE TRANSLATE · RULE-BASED

*"My heart has become garden garden."*

Word for word. The meaning is lost.

CHATGPT · GENERATIVE

"My heart is overjoyed — I'm blooming with happiness."

Meaning, culture, idiom — preserved.



— A BRIEF HISTORY

# How AI *evolved*, in four stages.

## 01 Rule-Based 1950S – 1980S

If-then logic. A vending machine: fixed input, fixed output.

## 02 Machine Learning 1990S – 2010S

Learns patterns from data. A spam filter that gets smarter over time.

## 03 Deep Learning 2010S

Neural networks. The AI that beat human champions at Go.

## 04 Foundation Models 2020 – NOW

Trained on the internet. Performs tasks it was never specifically trained for.

R  
*chef.*

the-Data / AI  
is a *recipe*  
*book.*

Generative  
AI is a *master*

10<sup>40</sup>

## 01 EXERCISE

*Seven minutes*

# Your first *generative output*.

*Open ChatGPT. Paste the prompt below, replacing the bracketed role with your own.*

## PROMPT

Write a 4-line poem about my job as a [your role].  
Make it warm and slightly funny.

## i Run it.

Hit enter.

## ii Pick a line.

The one that made you smile.

## iii Share it.

Paste it in chat

– II. THE LANDSCAPE –

# Key *milestones.*

*GPT, DALL·E, Gemini, Claude, Llama — what each one unlocked.*

— FIVE FAMILIES

# The *family tree* of modern AI.

— 2020 —

## GPT-3

Fluent language at scale.

— 2021 —

## DALL·E

Images from text.

— 2023 —

## GPT-4

Multimodal reasoning.

— 2023 —

## Gemini

Google ecosystem.

— 2023 / 24 —

## Claude · Llama

Safety & open source.

# 02

EXERCISE  
*Ten minutes*

## Model *speed dating*.

*Open ChatGPT, Claude, and Gemini in three tabs. Paste the same prompt into all three.*

### PROMPT

I'm a [marketing manager / finance analyst / doctor].  
Explain in 4 short bullet points why our quarterly revenue dipped 12%.  
Assume nothing about my business.  
Ask me 2 clarifying questions BEFORE answering.

OpenAI

ChatGPT

chat.openai.com

Anthropic

Claude

claude.ai

— FOR YOUR WORK

# What each model *unlocks*.

---

OPENAI

## ChatGPT

General drafting, brainstorming, ideation.

---

ANTHROPIC

## Claude

Long documents, nuanced analysis. The contract lawyer's tool.

---

GOOGLE

## Gemini

Lives inside Docs, Gmail, Sheets.

---

OPENAI · MIDJOURNEY

## DALL·E · Midjourney

Marketing visuals, mockups, concept art.

---

META

## Llama

Open source. Runs on private servers — banks, hospitals, government.

---

A HEURISTIC

*Pick the model for the job.*

Not the brand. Not the hype.

– III. MECHANICS –

# How *LLMs* work.

*Tokens, context windows, temperature.*



— TOKENS

# AI thinks in *LEGO bricks* of language.

— A FEW EXAMPLES —

strawberry

1

antidisestablishmentarianism

6

Shantanu Chandra

~5



1

*You're charged by the brick, not the sentence.*

*One thousand tokens ≈ seven hundred and fifty English words.*

*Hindi, Tamil and Arabic words can run three to five times more tokens than English.*

---

— **a** MICRO-EXERCISE  
*Four minutes*

---

See your *tokens* live.

Open [platform.openai.com/tokenizer](https://platform.openai.com/tokenizer)

— TYPE EACH, NOTE THE COUNT —

01 Your full name

---

02 A three-paragraph email you sent last week

---

03 An emoji of your choice

---

04 A Hindi or Tamil word in native script

---

THEN SHARE

What was your email's token count?

Type the number in chat.

18<sup>40</sup>

— CONTEXT WINDOW

# The model's *whiteboard*.

*Everything in the conversation lives on it. When it fills, the oldest content is erased.*

GPT-4O · OPENAI

128<sub>k</sub>

tokens · ~300 pages

CLAUDE SONNET · ANTHROPIC

200<sub>k</sub>

tokens · ~500 pages

GEMINI 1.5 PRO · GOOGLE


2<sub>m</sub>

tokens · ~5,000 pages

19<sup>/40</sup>

## — TEMPERATURE

# The *creativity dial*.



## FACTS

# "What is $2 + 2$ ?"

Always four. Temperature is irrelevant.

## CREATIVITY

"Name a coffee shop."

A different answer each time.

— **b** MICRO-EXERCISE  
*Three minutes*

Feel the *temperature*.

— NEW CHAT —

PROMPT

Give me 5 creative names for a wellness clinic in  
Mumbai.

— ANOTHER NEW CHAT —

PROMPT

Give me 5 creative names for a wellness clinic in  
Mumbai.

*Compare. The names are different.*

21<sup>/40</sup>

— B R E A K —

10:00

*Stand up. Water. Stretch.*

22<sup>/40</sup>

W elcome back.  
Type any  
*emoji* in chat.

– IV. THE STACK –

# The *value chain*.

*Models → platforms → applications → outcomes.*



— THE STACK

# Four *layers*, from brain to bottom line.

## 04 Outcomes — business value

Faster proposals · lower cost · fewer wasted hours

## 03 Applications — the tools you use

Microsoft Copilot · Notion AI · GitHub Copilot · ChatGPT

## 02 Platforms — secure infrastructure

Azure OpenAI · AWS Bedrock · Google Vertex AI · Hugging Face

## 01 Models — the raw brain

GPT-4 · Claude · Gemini · Llama

# 03

PAIRS · BREAKOUT

*Eight minutes*

## Reverse-engineer *the stack*.

*In your breakout pair, pick one tool. Figure out which layers it touches.*

MICROSOFT

### Copilot

AI in Word, Excel, Outlook.

NOTION

### Notion AI

AI inside documents.

GITHUB

### Copilot

AI that writes code.

BANKING

### WhatsApp bot

Customer support chatbot.

— V . R I S K —

# *Limitations.*

*Where AI breaks down — four hard limits.*

— FOUR HARD LIMITS

# Where a single LLM call *falls short*.

— HIGH RISK —

## Hallucination

Confident, wrong answers that look factual.

BITES AT: CITATIONS · NAMES · NUMBERS · DATES

— MEDIUM RISK —

## Static knowledge

Knows nothing after its training cutoff date.

BITES AT: CURRENT EVENTS · NEW POLICY · RECENT PRICES

— MEDIUM RISK —

## No memory

Forgets everything between conversations.

BITES AT: LONG PROJECTS · REPEAT CLIENTS · HANDOVERS

— LIMITING —

## No action

Can produce text. Cannot actually do things.

BITES AT: BOOKING · SENDING · SCHEDULING · CHECKING

— LIVE DEMO

# Confidently *wrong*.

PROMPT

Tell me about the 2019 Indian Supreme Court case 'Mehta vs. Sharma' on cryptocurrency regulation.  
Cite the bench and key judgment paragraphs.

THE TRICK

This case does not exist.

The model will write you a full judgment with fake judges, in confident detail.

— LIVE DEMO

# The model is *frozen in time*.

PROMPT

Who won the most recent Indian general election?  
What was the seat count?

FREE CHATGPT · NO WEB

*Stale, hedged. Sometimes confidently wrong.*

WITH WEB ACCESS

*Reaches out. Gets the real answer.*

30<sup>/40</sup>

## — LIVE DEMO

# Each conversation, a *fresh slate*.

## — CHAT ONE —

## PROMPT

My name is Priya. I'm a marketing director at an Indian D2C skincare brand. Remember this.

→ *"Got it, Priya."*

## — CHAT TWO · NEW WINDOW —

## PROMPT

What's my name and what do I do?

→ *"I don't have that information."*

— LIVE DEMO

It can *think*. It can *write*.  
It cannot *do*.

PROMPT

Book me a flight from Delhi to Bangalore for next Tuesday morning,  
add it to my Google Calendar,  
and email the itinerary to my assistant rajesh@company.com.

WHAT YOU GET

A detailed list of instructions.

It does not actually book, schedule, or send anything.



– VI. THE CLIMAX –

# The *hard ceiling* of chatbots.

*Where single LLM calls fall short — and what comes next.*

— THE SIX-STEP TEST

# One complex request. *Six failures.*

PROMPT

I run a small bakery in Mumbai. Help me launch a new product line:

- (1) research what's trending in the Mumbai bakery scene this month,
- (2) draft 3 Instagram posts,
- (3) generate 3 images for those posts,
- (4) schedule them across the next week,
- (5) email the schedule to sneha@bakery.com,
- (6) set a reminder to review performance in 30 days.

|   |    |                            |                          |
|---|----|----------------------------|--------------------------|
| ✗ | 01 | Research trending products | no live web on free tier |
| ✓ | 02 | Draft Instagram posts      |                          |
| ~ | 03 | Generate images            |                          |
| ✗ | 04 | Schedule posts             |                          |

34/40 works  
partial · tier limits  
cannot

A chatbot is a *brain  
in a jar*.  
It can think. It can  
write.

It cannot *browse ·  
remember · schedule · send ·  
check*.

— LIVE PREVIEW

# The *Agent Builder*.

*Same bakery task — with hands, eyes, and a calendar attached.*

## 01 ROLE System prompt

Plain English: "You are a marketing assistant..."

## 03 RUN Watch it work

Same task. Six steps. All completed.

## 02 TOOLS Integrations

Web search · calendar · email · memory · files.

04 FINAL EXERCISE  
Seven minutes

# Your *take-home prompt*.

*Pick one task you do at work every week. Write a prompt using this template.*

PROMPT

ROLE:            You are a [expert role].  
TASK:            I need you to [specific task].  
CONTEXT:        My situation is [brief background].  
CONSTRAINTS:   Keep it [tone, length, format].  
OUTPUT:         Give me [exact format you want].

DOCTOR

Discharge note

→ five-bullet shift handover.

MARKETING

Product launch

→ three Instagram caption variants.

## — RECAP

# What we *covered*.

## i Evolution

Recipe book → master chef.

## iii Mechanics

LEGO · whiteboard · dial.

## V Limitations

Hallucinate · static · no memory · no action.

## ii Milestones

Five model families.

## iv Value chain

Models → platforms → apps.

## vi Chatbot ceiling

Brain in a jar — no hands.

---

*Every failure has a fix. That fix is called an Agent.*

— COMING NEXT SESSION —

You'll *build* the Agent  
that succeeds at the bakery  
task.

*No code. Live.*

---

— HOMEWORK

# Take this *with you.*

*Reply to today's class email.*

---

— I —

One process at your job

... where GenAI could help today. Two sentences.

---

— II —

One place where it'd be dangerous

... where guardrails would be essential. Two sentences.

---

*Thank you. See you in Session 02.*

40<sup>/40</sup>